

# Continues Control (Single Agent)

I used DDPG algorithm to train the agent so agent can select best action against each state.

## Why DDPG:

DDPG is an actor-critic, model-free algorithm based on the deterministic policy gradient that can operate over continuous action spaces.

## Reference Used:

1. <https://arxiv.org/abs/1509.02971>
2. <https://arxiv.org/pdf/1509.02971.pdf>

## Summary of DDPG Network:

Used an Actor and Critic Network. As per the guidance of paper use Batch normalization as well. I used Adam optimizer to train the network. Fine tune the hyper parameters than mentioned in paper to converge faster.

- **Summary of Actor Network:**

Layer (type)	Output Shape	Param #
Linear-1	[-1, 128]	4,352
BatchNorm1d-2	[-1, 128]	256
Linear-3	[-1, 128]	16,512
Linear-4	[-1, 4]	516
=====		
Total params: 21,636		
Trainable params: 21,636		
Non-trainable params: 0		
=====		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.00		
Params size (MB): 0.08		
Estimated Total Size (MB): 0.09		
=====		

- **Summary of Critic Network:**

```
In [14]: summary(agent.critic_local, [(state_size,), (action_size, )])
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 128]	4,352
BatchNorm1d-2	[-1, 128]	256
Linear-3	[-1, 128]	17,024
Linear-4	[-1, 1]	129
=====		
Total params: 21,761		
Trainable params: 21,761		
Non-trainable params: 0		
=====		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.00		
Params size (MB): 0.08		
Estimated Total Size (MB): 0.09		
=====		

## Summary Hyper parameter:

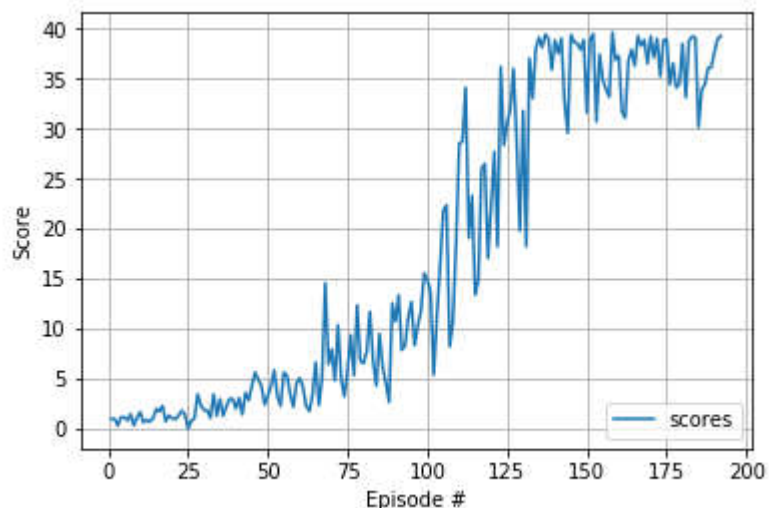
- **Maximum steps per episode:** 1000
- **Replay Buffer size:** 100000
- **Batch Size:** 128
- **Gamma:** 0.99
- **Tau:** 0.001
- **Actor Learning Rate:** 0.0002
- **Critic Learning Rate:** 0.0002
- **Weight Decay:** 0.0

## Rewards Function Performance During Training:

- **Average Score after every 100 episodes:**

```
Episode 100      Average Score: 4.51
Episode 192      Average Score: 30.19
Environment solved in 92 episodes!      Average Score: 30.19
```

- **Plot shows average rewards against each episode:**



## Rewards Function Performance During Prediction:

```
Episode 98      Average Score: 35.21
Episode 99      Average Score: 35.18
Episode 100     Average Score: 35.21
Average Score Of 100 Consecutive Episodes: 35.13582271759092
```

## Networks want to try in future:

1. D4PG for multiagent.
2. Proximal Policy Optimization.